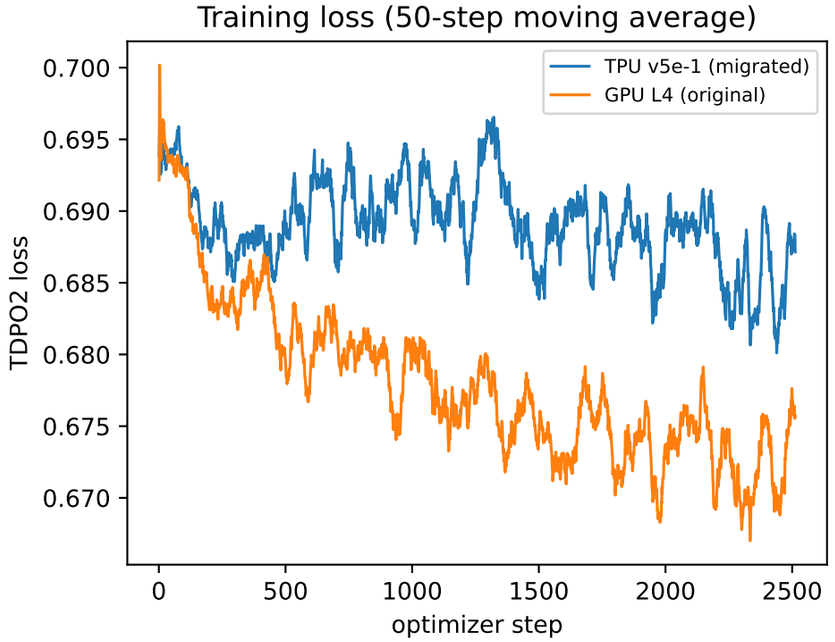
Training loss (50-step moving average)
TDPO2 loss
optimizer step
TPU v5e-1 (migrated)
GPU L4 (original)

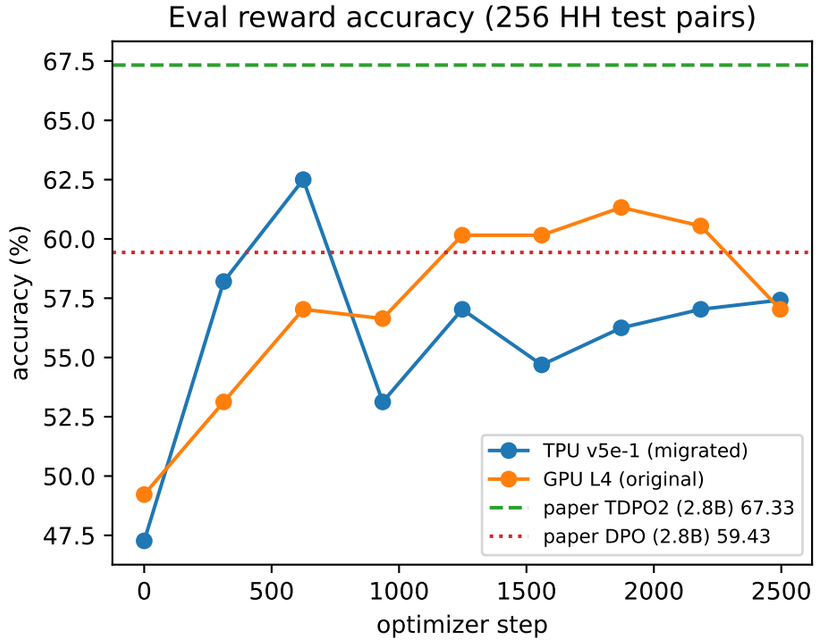
Eval reward accuracy (256 HH test pairs)
accuracy (%)
optimizer step
TPU v5e-1 (migrated)
GPU L4 (original)
paper TDPO2 (2.8B) 67.33
paper DPO (2.8B) 59.43

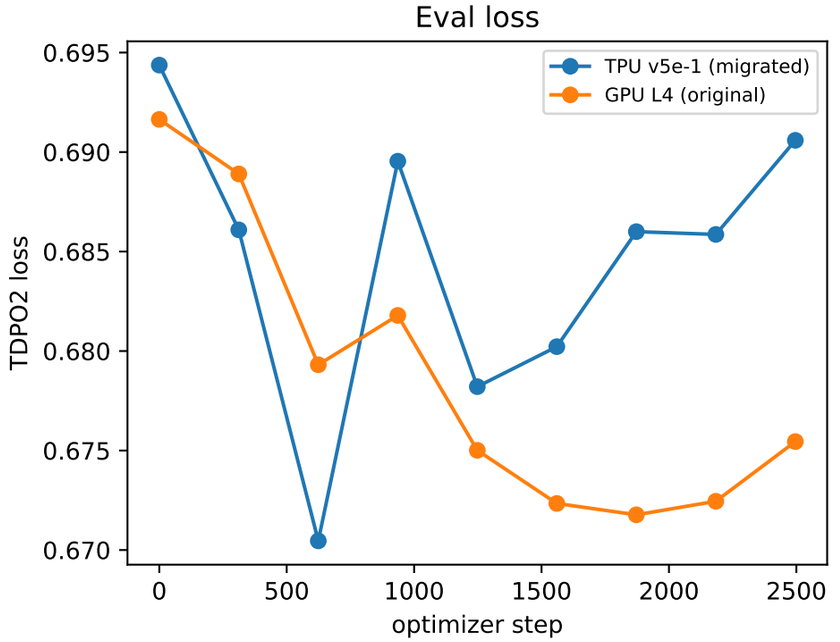
Eval loss
TDPO2 loss
optimizer step
TPU v5e-1 (migrated)
GPU L4 (original)