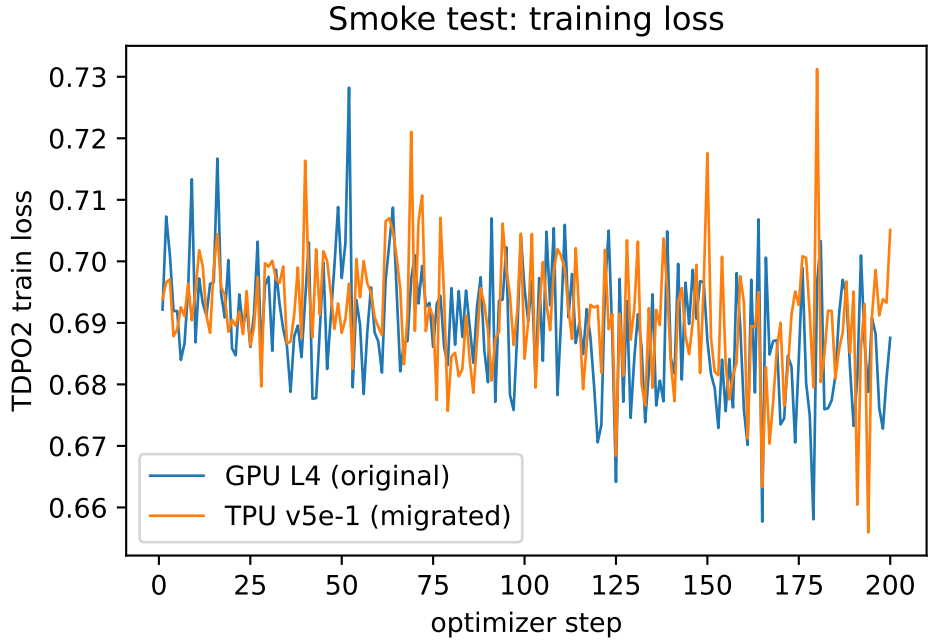
Smoke test: training loss
TDPO2 train loss
optimizer step
GPU L4 (original)
TPU v5e-1 (migrated)
0.73
0.72
0.71
0.70
0.69
0.68
0.67
0.66
0
25
50
75
100
125
150
175
200

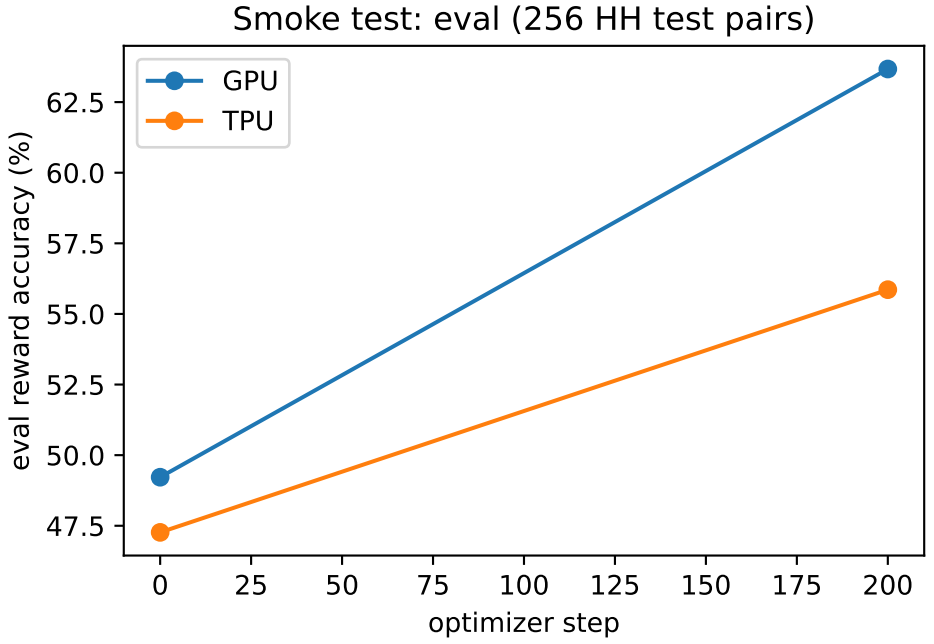
Smoke test: eval (256 HH test pairs)
eval reward accuracy (%)
optimizer step
GPU
TPU
62.5
60.0
57.5
55.0
52.5
50.0
47.5
0
25
50
75
100
125
150
175
200